Demo Use case 1 - SNOWFLAKE.CORTEX.SUMMARIZE LLM function

This lab walks you through a hands-on exercise using the Snowflake Cortex LLM function SNOWFLAKE.CORTEX.SUMMARIZE to create concise summaries from support transcripts. It includes a sample dataset, step-by-step commands, expected outputs, validation checks, troubleshooting tips, and extension tasks you can try.

Use this in a development Snowflake account or sandbox. The lab assumes you have a role with privileges to create warehouses, databases, schemas, tables and to execute Cortex functions.

Learning objectives

1. Load a small set of text documents (support transcripts) into Snowflake.
2. Call SNOWFLAKE.CORTEX.SUMMARIZE to produce concise summaries.
3. Inspect, evaluate, and refine summarization results.
4. Build a simple table of summaries for use in downstream workflows (agent notes, dashboards).
5. Practice troubleshooting common issues.

Environment setup (one-time)

Run these commands in Snowsight or any Snowflake SQL client.

```
-- Choose an account role that can create objects
USE ROLE ACCOUNTADMIN;
-- Create a warehouse for the lab (small size for cost control)
CREATE OR REPLACE WAREHOUSE lab_wh WITH WAREHOUSE_SIZE =
'XSMALL' AUTO_SUSPEND = 60 AUTO_RESUME = TRUE;


-- Create a database and schema
CREATE OR REPLACE DATABASE lab_db;
CREATE OR REPLACE SCHEMA lab_db.public;


USE WAREHOUSE lab_wh;
USE DATABASE lab_db;
USE SCHEMA lab_db.public;
```

Step 1 — Create sample dataset

Create a table of customer support transcripts and populate it with example rows.

```
CREATE OR REPLACE TABLE customer_support_tickets (
  ticket_id INTEGER,
  created_at TIMESTAMP_NTZ,
  customer_name VARCHAR,
  language VARCHAR,
  category VARCHAR,
  transcript STRING
);

INSERT INTO customer_support_tickets (ticket_id, created_at, customer_name,
language, category, transcript) VALUES
(1001, CURRENT_TIMESTAMP(), 'Alice Brown', 'English', 'Delivery',
 'The product was delayed by three days. I did not receive any shipping updates. I tried
calling support twice and email but no response. I need it for a birthday.'),
(1002, CURRENT_TIMESTAMP(), 'Rajesh Kumar', 'Hindi', 'Payment',
 'मेरा भुगतान विफल हो गया लेकिन मेरे खाते से पैसे कट गए। मैंने बैंक से भी पुष्टि कर ली है। कृपया पैसे
वापस भेजें।'),
(1003, CURRENT_TIMESTAMP(), 'Maria Lopez', 'Spanish', 'Product',
 'El producto llegó con la pantalla rota. Necesito un reemplazo o reembolso. Adjunto
fotos.'),
```

```
(1004, CURRENT_TIMESTAMP(), 'John Smith', 'English', 'Refund',
 'I requested a refund two weeks ago and the refund has not been processed yet. The order number is ORD-54321. Can you confirm the status?'),
(1005, CURRENT_TIMESTAMP(), 'Sophie Dubois', 'French', 'Quality',
 'Le produit ne correspond pas à la description sur le site. Les dimensions sont incorrectes et la finition est mauvaise.');
```

Step 2 — Basic summarization (single column)

Run a simple summarization query to produce one-line summaries for every transcript.

```
SELECT
  ticket_id,
  customer_name,
  language,
  category,
  SNOWFLAKE.CORTEX.SUMMARIZE(transcript) AS auto_summary
FROM customer_support_tickets
ORDER BY ticket_id;
```

What to expect

A result set with an auto_summary column where each transcript has a concise summary. Example expected outputs (samples, wording may vary):

ticket_id	auto_summary
1001	Customer reports a three day delivery delay and lack of shipping updates; attempted contact with support.
1002	Payment failed but money was deducted from account; customer requests refund.
1003	Product arrived with a broken screen; customer requests replacement or refund.
1004	Refund requested two weeks ago for order ORD-54321; customer requests status update.
1005	Product does not match website description; incorrect dimensions and poor finish.

Note: Exact phrasing can vary because the summarization model may produce slightly different wording.

Step 3 — Summarization with context columns

Sometimes you want the summary to include metadata, for example language or ticket id. Use CONCAT to include additional context in the input text so the model can incorporate it into the summary.

```
SELECT
  ticket_id,
  SNOWFLAKE.CORTEX.SUMMARIZE(
    CONCAT('Ticket ', ticket_id, ' [', language, '] Transcript: ', transcript)
  ) AS summary_with_context
FROM customer_support_tickets
```

```
ORDER BY ticket_id;
```

Why do this

Including small structured context helps produce summaries that mention the ticket id or language when desirable, useful for auditability or downstream automation.

Step 4 — Shorter or longer summaries (technique)

SNOWFLAKE.CORTEX.SUMMARIZE usually returns concise text by default. If you need a stricter length target or a specific format (for example, a one-line bullet or JSON), combine COMPLETE (a general LLM completion function) with an instruction prompt, or post-process the summary. Example of constrained output using COMPLETE:

```
SELECT
 ticket_id,
 SNOWFLAKE.CORTEX.COMPLETE(
  'mistral-large',
  CONCAT(
   'Produce a one-sentence summary in English for the following transcript. ',
   'Return only the sentence. Transcript: ', transcript
  )
 ) AS one_sentence_summary
FROM customer_support_tickets;
```

Note: COMPLETE gives more control over instructions and format. Use it when you must enforce exact output structure.

Step 5 — Write summaries back to a table for reporting

Store generated summaries in a new table for downstream reports, dashboards, or agent interfaces.

```
CREATE OR REPLACE TABLE ticket_summaries AS
SELECT
  ticket_id,
  customer_name,
  language,
  category,
  SNOWFLAKE.CORTEX.SUMMARIZE(transcript) AS summary,
  CURRENT_TIMESTAMP() AS summary_generated_at
FROM customer_support_tickets;
```

Validate inserted rows:

```
SELECT * FROM ticket_summaries ORDER BY ticket_id;
```

Step 6 — Evaluation and quality checks

Create a simple quality framework to check summarization outputs.

1. Sample verification

- Manually inspect 5 to 10 random summaries for faithfulness and usefulness.

2. Automated checks

- Check for empty summaries

```
SELECT COUNT(*) AS empty_count
FROM ticket_summaries
WHERE TRIM(summary) = '';
```

- Check for summaries that repeat the whole input (indicating poor summarization)

```
SELECT ticket_id
FROM ticket_summaries s
JOIN customer_support_tickets t USING (ticket_id)
WHERE LENGTH(s.summary) > LENGTH(t.transcript) * 0.8;
```

3. Accuracy sampling

- Create a small human-labelled dataset (ground truth) of 50 transcripts with reference summaries and compute simple overlap metrics (ROUGE-like or manual scoring).

Suggested manual scoring rubric (for human reviewers)

- Score 0: Incorrect or misleading summary

- Score 1: Partial, misses critical point
- Score 2: Accurate and concise

Log results in a table for tracking model quality over time.