

# Core Concepts of Large Language Models (LLMs)

---

## 1. Introduction

Large Language Models (LLMs) such as GPT, Claude, and Gemini rely on a set of **core configuration concepts and parameters** that directly influence how responses are generated. Understanding these concepts is essential for **prompt engineering, application development, API integration, and responsible AI usage**.

This document explains the most important core concepts including **tokens, temperature, system roles, and other key parameters**, with examples and real-world software development use cases.

---

## 2. Tokens: The Fundamental Unit of LLM Processing

### 2.1 What Are Tokens?

A **token** is a chunk of text that an LLM processes. Tokens are **not the same as words**.

- A word may be split into multiple tokens
- Punctuation and spaces can be tokens
- Numbers and symbols are also tokenized

### Example Tokenization

Text	Tokens (Approximate)
Hello	1 token
ChatGPT	1 token
ChatGPT is powerful	4–5 tokens
Internationalization	3–4 tokens

### 2.2 Why Tokens Matter

Tokens impact:

- Cost (API pricing is token-based)
- Context window size
- Performance and latency
- Prompt and response length limits

## **2.3 Input Tokens vs Output Tokens**

<b>Token Type</b>	<b>Description</b>
Input tokens	Tokens sent in the prompt
Output tokens	Tokens generated by the model
Total tokens	Input + Output

### **Example**

Prompt: Explain REST API (10 tokens)

Response: Detailed explanation (120 tokens)

Total usage: 130 tokens

---

## **2.4 Use Case: Token Awareness in Software Development**

### **Scenario:**

A chatbot integrated into a customer support system must handle long conversations.

### **Solution:**

- Limit conversation history
  - Summarize older messages
  - Control max output tokens
-

### **3. Temperature: Controlling Creativity and Randomness**

#### **3.1 What Is Temperature?**

**Temperature** controls how **random or deterministic** the model's output is.

Temperature Value	Behavior
0.0 – 0.2	Very deterministic, factual
0.3 – 0.5	Balanced
0.6 – 0.8	Creative
0.9 – 1.0	Highly creative, less predictable

---

#### **3.2 Example**

##### **Prompt:**

Generate a product description for a smartwatch

Temperature	Output Style
0.2	Technical and precise
0.7	Marketing-oriented and engaging
1.0	Very creative but may exaggerate

### 3.3 Best Practices

Use Case	Recommended Temperature
Code generation	0.0 – 0.3
Chatbots	0.4 – 0.6
Creative writing	0.7 – 0.9

---

## 4. System, User, and Assistant Roles

### 4.1 Role-Based Prompting

LLMs support **role-based messages** to control behavior.

Role	Purpose
System	Sets rules, tone, and behavior
User	Provides instructions or questions
Assistant	Model-generated response

---

### 4.2 System Role (Most Important)

The **system role** defines how the model should behave.

#### Example

System: You are a strict JSON API. Return only valid JSON.

User: Generate customer data.

Result:

- Output strictly follows structure
- No explanations added

---

### 4.3 Use Case: Enterprise Chatbot

## **System role:**

You are a banking assistant.

Never provide financial advice.

Respond formally.

This ensures **compliance and consistency**.

---

## **5. Max Tokens: Limiting Response Length**

### **5.1 What Is Max Tokens?**

Defines the **maximum number of tokens** the model can generate in a response.

---

### **5.2 Why It Is Important**

- Prevents long, unnecessary responses
  - Controls cost
  - Improves performance
- 

### **5.3 Example**

<b>Max Tokens</b>	<b>Output</b>
50	Short summary
300	Detailed explanation
1000	Long technical document

---

## 5.4 Use Case

### Scenario:

Generating API error messages.

### Solution:

Set max\_tokens = 50 to avoid verbose outputs.

---

## 6. Top-p (Nucleus Sampling)

### 6.1 What Is Top-p?

Top-p limits token selection to the **most probable tokens whose cumulative probability  $\leq p$** .

Top-p Value	Effect
0.1	Very conservative
0.5	Balanced
0.9	More diverse

---

### 6.2 Temperature vs Top-p

Parameter	Controls
Temperature	Randomness
Top-p	Probability mass

### Best Practice:

Use **either temperature or top-p**, not both aggressively.

---

## 7. Frequency and Presence Penalties

### 7.1 Frequency Penalty

- Reduces repeated words or phrases

### 7.2 Presence Penalty

- Encourages introducing new topics
- 

### 7.3 Example

#### Without penalty:

This API is fast. This API is reliable. This API is scalable.

#### With penalty:

This API is fast, reliable, and scalable.

---

## 8. Practical Use Case: API-Based AI Assistant

### Scenario

An organization builds an AI assistant for developers.

### Configuration

Parameter	Value
System role	“You are a senior software architect”
Temperature	0.3
Max tokens	300
Top-p	0.9
Frequency penalty	0.2

## **Result**

- Accurate
  - Concise
  - Professional responses
- 

## **9. Common Mistakes to Avoid**

- Ignoring token limits
  - Using high temperature for code generation
  - Not defining system role
  - Mixing explanation with structured output
  - Overusing long prompts unnecessarily
- 

## **10. Summary Table of Core Parameters**

<b>Concept</b>	<b>Purpose</b>	<b>Typical Use</b>
Tokens	Measure text size	Cost & limits
Temperature	Controls creativity	Style control
System role	Sets behavior	Governance
Max tokens	Output length	Cost control
Top-p	Probability filtering	Diversity
Penalties	Reduce repetition	Clean output

---

## 11. Conclusion

Understanding core LLM concepts such as **tokens, temperature, roles, and generation parameters** is critical for building reliable, cost-effective, and production-grade AI applications. These parameters allow developers to control **accuracy, creativity, safety, and consistency**, making LLMs suitable for real-world software systems.

When applied correctly, these concepts transform LLMs from simple chat tools into **powerful, configurable AI engines**.